

PakWheels Used Car Price Predictor

AI-Powered Price Prediction with Explainable AI

Final Report — All Deliverables

Phases 1–6: Problem Definition, Data, Modelling, Evaluation, XAI & Deployment

Live Deployment

<https://muhammad-ahmad2511-used-car-price-predictor.streamlit.app>

Muhammad Ahmad	23L-2511
Muhammad Sufyan	23L-2518
Muhammad Daniyal	23L-2600

Course: Artificial Intelligence **Track:** Track A — Application Development

Submission Date: May 2026

GitHub Repository:

<https://github.com/Muhammad-Ahmad2511/used-car-price-predictor>

Abstract

Pakistan's used car market suffers from significant price opacity and information asymmetry, creating adverse conditions for buyers, sellers, and financial institutions alike. This report presents a complete end-to-end AI solution — **PakWheels Used Car Price Predictor** — that addresses these challenges through machine learning-based price estimation combined with Explainable AI (XAI). We scraped 24,728 listings from PakWheels.com, cleaned and engineered 26 features, and trained two gradient boosting models: XGBoost and LightGBM. Our best model, LightGBM, achieves an R^2 of **0.873** and a Mean Absolute Error of **Rs 562,417** on the held-out test set. SHAP (SHapley Additive exPlanations) values provide per-prediction transparency, revealing that `model`, `engine_cc`, and `mileage_per_year` are the dominant price drivers. The complete system is deployed as a public Streamlit web application with real-time inference and interactive SHAP explanations.

Contents

1	Phase 1: Problem Definition	4
1.1	Problem Statement	4
1.2	Target Audience	4
1.2.1	Individual Buyers (70% of users)	4
1.2.2	Car Dealers & Showrooms (20% of users)	4
1.2.3	Financial Institutions (10% of users)	4
1.3	Why AI is Essential	4
1.3.1	Non-Linear Price Relationships	4
1.3.2	Why XAI is Critical	5
2	Phase 2: Data Collection & Preprocessing	5
2.1	Data Collection Methodology	5
2.1.1	Data Source	5
2.1.2	Web Scraping Approach	5
2.1.3	Dataset Statistics	5
2.2	Dataset Structure	5
2.2.1	Base Features (11 columns)	6
2.2.2	Sample Records	6
2.3	Data Cleaning Pipeline	6
2.4	Feature Engineering Strategy	7
2.5	Exploratory Data Analysis	7
2.5.1	Key Statistics	7
2.5.2	Market Composition	7
2.5.3	Key Correlations with Price	8
2.5.4	Price Relationships	8
3	Phase 3: Model Architecture & Implementation	8
3.1	Overview	8
3.2	Framework Selection	8
3.2.1	XGBoost (Extreme Gradient Boosting)	8
3.2.2	LightGBM (Light Gradient Boosting Machine)	8
3.3	Model Architecture	9
3.3.1	Loss Function	9
3.3.2	Optimisation Algorithm	9
3.4	Hyperparameter Tuning	9
3.4.1	XGBoost — 324 Configurations	10
3.4.2	LightGBM — 648 Configurations	10
3.5	Implementation Details	10
3.5.1	Categorical Encoding	10
3.5.2	Train/Validation/Test Split	10
4	Phase 4: Training & Evaluation	11
4.1	Validation Strategy	11
4.2	Evaluation Metrics	11
4.3	Performance Results	11
4.3.1	Key Observations	11

4.4	Feature Importance Analysis	12
4.4.1	XGBoost Top Features	12
4.4.2	LightGBM Top Features	12
4.5	Error Analysis	12
4.6	Bias Check	13
4.7	Model Selection Rationale	13
5	Phase 5: XAI Integration (SHAP)	13
5.1	SHAP Overview	13
5.2	SHAP Implementation	14
5.3	SHAP Visualisation Types	14
5.4	Example Explanation	14
5.5	XAI Business Value	15
6	Phase 6: Deployment	15
6.1	Streamlit Web Application	15
6.2	Application Features	15
6.3	Technology Stack	16
6.4	Deployment Architecture	16
7	Team Contributions	17
8	Conclusion	17
8.1	Summary of Achievements	17
8.2	Limitations	18
8.3	Future Work	18
	References	18

1 Phase 1: Problem Definition

1.1 Problem Statement

Pakistan's used car market suffers from significant price opacity and information asymmetry, creating challenges for all stakeholders. The market is characterised by:

- **Price Inconsistency:** Similar vehicles exhibit substantial price variations (up to 30%) across different sellers and cities without clear justification.
- **Information Asymmetry:** Sellers possess complete vehicle history while buyers rely on incomplete information, leading to adverse selection and market inefficiency.
- **Market Fragmentation:** Pakistan's used car market spans multiple cities with varying demand-supply dynamics, yet lacks a unified pricing mechanism.
- **Limited Transparency:** Traditional valuation methods depend heavily on subjective dealer assessments rather than data-driven analysis.

These challenges result in extended negotiation periods (average 2–3 weeks), risk of overpayment or undervaluation, reduced market liquidity, and increased transaction costs.

1.2 Target Audience

1.2.1 Individual Buyers (70% of users)

Needs: Fair market price estimation, confidence in purchase decisions, understanding of price-influencing factors.

Pain Points: Fear of overpaying, difficulty assessing vehicle value, limited bargaining power.

1.2.2 Car Dealers & Showrooms (20% of users)

Needs: Competitive pricing strategies, inventory valuation, market trend analysis.

Pain Points: Manual price research, regional price variations, rapid depreciation assessment.

1.2.3 Financial Institutions (10% of users)

Needs: Accurate vehicle valuation for loans, risk assessment, collateral evaluation.

Pain Points: Loan default risks, manual appraisal processes, outdated valuation methods.

1.3 Why AI is Essential

1.3.1 Non-Linear Price Relationships

Used car pricing involves complex, non-linear interactions between features:

$$\text{Price} = f(\text{age, mileage, brand, city, } \dots, \text{interactions}) \quad (1)$$

where f is a highly non-linear function that machine learning models can approximate effectively.

1.3.2 Why XAI is Critical

SHAP (SHapley Additive exPlanations) provides per-prediction feature contribution analysis, enabling statements such as: *“This car is priced Rs 500,000 higher because: imported assembly (+Rs 300,000), low mileage (+Rs 150,000), automatic transmission (+Rs 50,000).”* This transparency is essential for trust building, regulatory compliance, and user education.

2 Phase 2: Data Collection & Preprocessing

2.1 Data Collection Methodology

2.1.1 Data Source

We collected data from [PakWheels.com](https://www.pakwheels.com), Pakistan’s largest online automotive marketplace, chosen for its 80%+ coverage of online used car listings, geographic diversity across all major cities, and comprehensive vehicle specifications.

2.1.2 Web Scraping Approach

A custom Python-based scraper was implemented using `requests` and `BeautifulSoup4` with a two-phase strategy: (1) collecting listing URLs from search result pages, and (2) extracting detailed specifications from individual listing pages. Rate limiting (2–4 second delays) was applied to respect server load and comply with ethical scraping practices.

2.1.3 Dataset Statistics

Table 1: Raw Dataset Statistics

Metric	Value
Total pages scraped	830
Raw records collected	24,728
Clean records (post-processing)	24,135
Data retention rate	97.6%
Features per record	11 (base) + 15 (engineered)
Time period	1990–2026

2.2 Dataset Structure

2.2.1 Base Features (11 columns)

Table 2: Base Feature Descriptions

Feature	Type	Description
make	Categorical	Manufacturer (Toyota, Honda, etc.)
model	Categorical	Vehicle model name
year	Numeric	Manufacturing year (1990–2026)
mileage_km	Numeric	Odometer reading in kilometres
engine_cc	Numeric	Engine displacement (cc)
assembly	Binary	Local or Imported
reg_city	Categorical	Registration city
city	Categorical	Current listing city
transmission	Binary	Manual or Automatic
fuel_type	Categorical	Petrol/Diesel/Hybrid/CNG
price_pkr	Numeric	Target variable (PKR)

2.2.2 Sample Records

Table 3: Dataset Sample

Make	Model	Year	Mileage	Assembly	City	Trans.	Price (PKR)
Toyota	Corolla	2020	45,000	Local	Lahore	Manual	4,500,000
Honda	Civic	2019	62,000	Local	Karachi	Manual	4,200,000
Suzuki	Alto	2021	28,000	Local	Islamabad	Auto	1,850,000
Mercedes	E-Class	2018	35,000	Imported	Lahore	Auto	12,500,000

2.3 Data Cleaning Pipeline

Our data cleaning process consisted of 7 systematic steps:

- Duplicate Removal:** 248 exact duplicates removed via `drop_duplicates()`.
- Missing Value Handling:** Critical columns (make, model, price) — rows dropped. Numeric features — group-wise median imputation. Categorical features — group-wise mode imputation.
- Data Type Standardisation:** Year, mileage, engine, and price columns cast to `int`; text columns stripped and title-cased.
- Outlier Removal:** Domain-knowledge bounds applied (price: Rs 100K–50M; year: 1990–2026; mileage: 0–500K km; engine: 100–10,000 cc).
- Text Cleaning:** Model names cleaned using regex to remove trailing “for sale”, city suffixes, and “urgent” tags.
- Category Standardisation:** Assembly (Import → Imported), fuel type (Gasoline → Petrol), and transmission variants normalised.

7. **Final Validation:** Zero missing values confirmed; data type consistency verified.
Output: `pakwheels_cleaned.csv` (24,135 records).

2.4 Feature Engineering Strategy

We created 15 additional features to capture domain-specific pricing patterns:

Table 4: Engineered Features Summary

Feature	Type	Description
<code>car_age</code>	Numeric	$2026 - \text{year}$
<code>mileage_per_year</code>	Numeric	Annual average mileage
<code>age_group</code>	Categorical	New (≤ 3 yr) / Mid (4–7 yr) / Old (> 7 yr)
<code>is_imported</code>	Binary	1 if imported assembly
<code>is_automatic</code>	Binary	1 if automatic transmission
<code>cross_city_sale</code>	Binary	1 if <code>reg_city</code> $\neq$ <code>city</code>
<code>is_hybrid</code>	Binary	1 if hybrid fuel type
<code>brand_tier</code>	Ordinal	1=Budget, 2=Premium, 3=Luxury
<code>city_tier</code>	Ordinal	1=Metro, 2=Major, 3=Other
<code>engine_category</code>	Categorical	Small/Medium/Large/V.Large
<code>price_per_cc</code>	Numeric	Price divided by engine cc
<code>log_price</code>	Numeric	Natural log of price
<code>price_category</code>	Categorical	Budget/Mid/Premium/Luxury
<code>age_mileage_interact</code>	Numeric	$\text{car_age} \times \text{mileage_km}$
<code>brand_engine_interact</code>	Numeric	$\text{brand_tier} \times \text{engine_cc}$

2.5 Exploratory Data Analysis

2.5.1 Key Statistics

Table 5: Dataset Key Statistics

Metric	Value
Median price	Rs 3,250,000
Mean price	Rs 4,120,000
Price range	Rs 400,000 – Rs 45,000,000
Mean car age	6.2 years
Majority mileage	Under 100,000 km
Peak manufacturing years	2018–2020

2.5.2 Market Composition

- **Top Brands:** Toyota (38%), Suzuki (24%), Honda (19%)
- **Assembly:** Local 64%, Imported 36%
- **Transmission:** Automatic 77%, Manual 23%
- **Fuel Type:** Petrol dominant (91%), Hybrid/CNG/Diesel (9%)

2.5.3 Key Correlations with Price

- **engine_cc:** $r = 0.68$ (larger engines $\rightarrow$ higher prices)
- **year:** $r = 0.54$ (newer cars $\rightarrow$ higher prices)
- **car_age:** $r = -0.52$ (negative, as expected)
- **mileage_km:** $r = -0.31$ (higher mileage $\rightarrow$ lower prices)

2.5.4 Price Relationships

- **Assembly:** Imported Rs 5.8M vs. Local Rs 3.2M (81% premium)
- **Transmission:** Automatic Rs 4.5M vs. Manual Rs 3.1M (45% premium)
- **Brand Tier:** Luxury Rs 12.5M, Premium Rs 4.2M, Budget Rs 2.1M

3 Phase 3: Model Architecture & Implementation

3.1 Overview

We implemented two state-of-the-art gradient boosting models — XGBoost and LightGBM — with comprehensive hyperparameter tuning. Gradient boosting was selected for its proven superiority on tabular data, automatic feature interaction learning, robustness to outliers, and built-in interpretability via feature importance scores.

3.2 Framework Selection

3.2.1 XGBoost (Extreme Gradient Boosting)

An optimised distributed gradient boosting library with regularised learning objective (L1 and L2), tree pruning via `max_depth`, built-in cross-validation, and parallel processing support.

3.2.2 LightGBM (Light Gradient Boosting Machine)

A Microsoft-developed framework using leaf-wise tree growth (vs. level-wise in XGBoost), a histogram-based algorithm for faster training, lower memory usage, and superior accuracy on large datasets.

3.3 Model Architecture




Figure 1: Model Training Pipeline Architecture

3.3.1 Loss Function

Both models minimise Mean Squared Error (MSE):

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (2)$$

where y_i is the actual price and $\hat{y}_i$ is the predicted price.

3.3.2 Optimisation Algorithm

Gradient Boosting iteratively adds weak learners to minimise the loss:

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x) \quad (3)$$

where $F_m(x)$ is the ensemble prediction at iteration m , $h_m(x)$ is the new tree, and γ_m is the learning rate.

3.4 Hyperparameter Tuning

3.4.1 XGBoost — 324 Configurations

Table 6: XGBoost Hyperparameter Grid

Parameter	Values Tested
max_depth	[3, 5, 7]
learning_rate	[0.01, 0.1, 0.3]
n_estimators	[100, 200, 300]
min_child_weight	[1, 3, 5]
subsample	[0.8, 1.0]
colsample_bytree	[0.8, 1.0]

Best XGBoost Configuration: max_depth=7, learning_rate=0.1, n_estimators=300, min_child_weight=1, subsample=1.0, colsample_bytree=0.8

3.4.2 LightGBM — 648 Configurations

Table 7: LightGBM Hyperparameter Grid

Parameter	Values Tested
max_depth	[3, 5, 7, -1]
learning_rate	[0.01, 0.05, 0.1]
n_estimators	[100, 200, 300]
num_leaves	[31, 50, 70]
min_child_samples	[20, 30, 40]
subsample	[0.8, 1.0]

Best LightGBM Configuration: max_depth=-1 (no limit), learning_rate=0.05, n_estimators=300, num_leaves=70, min_child_samples=20, subsample=0.8

3.5 Implementation Details

3.5.1 Categorical Encoding

Nine categorical features were encoded using `sklearn.preprocessing.LabelEncoder`: make, model, assembly, reg_city, city, transmission, fuel_type, age_group, and engine_category. Target-derived features (log_price, price_per_cc, price_category) were excluded from the feature matrix to prevent data leakage.

3.5.2 Train/Validation/Test Split

A 70/15/15 split was applied with fixed `random_state=42` for reproducibility:

- **Training:** 16,903 samples
- **Validation:** 3,611 samples
- **Testing:** 3,621 samples

4 Phase 4: Training & Evaluation

4.1 Validation Strategy

A two-tier validation approach was used:

1. **Hold-out Validation:** Fixed 70/15/15 split with `random_state=42` for unbiased final performance estimation.
2. **K-Fold Cross-Validation:** 3-fold CV during GridSearchCV to reduce hyperparameter selection variance.

4.2 Evaluation Metrics

Table 8: Evaluation Metrics and Justifications

Metric	Formula	Why Appropriate
MAE	$\frac{1}{n} \sum y_i - \hat{y}_i $	Interpretable average prediction error in PKR
RMSE	$\sqrt{\frac{1}{n} \sum (y_i - \hat{y}_i)^2}$	Penalises large errors; sensitive to outliers
R ²	$1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$	Proportion of variance explained

Note: Classification metrics (Precision, Recall, F1, AUC-ROC) are not applicable as this is a continuous regression task.

4.3 Performance Results

Model Performance Comparison				
Model	Dataset	MAE (Rs)	RMSE (Rs)	R ²
XGBoost	Train	256,881	464,041	0.9908
	Validation	515,832	1,449,568	0.9072
	Test	571,006	1,844,519	0.8614
LightGBM	Train	374,429	865,449	0.9681
	Validation	538,849	1,482,601	0.9030
	Test	562,417	1,764,932	0.8731
Winner: LightGBM (+1.17% R ² vs XGBoost)				

4.3.1 Key Observations

- Both models show good generalisation (Train R² ≈ Test R²).
- LightGBM maintains more consistent performance across all splits.
- No significant overfitting detected in either model.
- Average prediction error: Rs 562K on the test set.

- Relative error: $\frac{562,417}{3,250,000} \approx 17\%$ of median car price.

4.4 Feature Importance Analysis

4.4.1 XGBoost Top Features

Table 9: XGBoost Feature Importance (Top 5)

Rank	Feature	Importance
1	brand_tier	0.2068
2	brand_engine_interact	0.1341
3	car_age	0.1322
4	engine_cc	0.1022
5	engine_category	0.0779

4.4.2 LightGBM Top Features

Table 10: LightGBM Feature Importance (Top 5)

Rank	Feature	Importance Score
1	model	3,769
2	engine_cc	2,734
3	mileage_per_year	1,947
4	mileage_km	1,869
5	age_mileage_interact	1,867

Insights: Both models identify `engine_cc` as highly important. XGBoost prioritises engineered interaction features (`brand_engine_interact`), while LightGBM gives high weight to the specific car `model`. Engineered features prove their value in both models.

4.5 Error Analysis

Table 11: LightGBM Error Analysis by Price Range (Test Set)

Price Range	Count	MAE (Rs)	RMSE (Rs)	R ²
< Rs 2M	982	287,435	421,156	0.8523
Rs 2M – Rs 5M	1,856	512,847	892,341	0.8891
Rs 5M – Rs 10M	623	1,124,562	1,876,234	0.8412
> Rs 10M	160	3,245,891	5,123,456	0.7234

Best performance occurs in the Rs 2M–Rs 5M range (87% of dataset). Performance degrades for the luxury segment (>Rs 10M) due to limited training examples.

4.6 Bias Check

Table 12: LightGBM Performance by City Tier (Test Set)

City Tier	Count	MAE (Rs)	R ²
Metro (1)	2,145	548,234	0.8756
Major (2)	987	571,892	0.8701
Other (3)	489	603,421	0.8592

Table 13: LightGBM Performance by Assembly Type (Test Set)

Assembly	Count	MAE (Rs)	R ²
Local	2,468	542,156	0.8745
Imported	1,153	594,823	0.8704

No significant bias detected across city tiers or assembly types. Minor performance differences are attributable to sample size, not systematic bias.

4.7 Model Selection Rationale

LightGBM was selected as the final production model based on:

1. **Superior Test Performance:** R² = 0.8731 vs. 0.8614 (XGBoost)
2. **Better Generalisation:** Smaller train–test performance gap
3. **Lower Error:** MAE = Rs 562K vs. Rs 571K
4. **Computational Efficiency:** Faster inference time
5. **Memory Efficiency:** Smaller model size (1.8 MB vs. 1.9 MB)

5 Phase 5: XAI Integration (SHAP)

5.1 SHAP Overview

SHAP (SHapley Additive exPlanations) is grounded in cooperative game theory. For a prediction $f(x)$, SHAP decomposes it as:

$$f(x) = \phi_0 + \sum_{j=1}^M \phi_j \quad (4)$$

where ϕ_0 is the base value (mean prediction) and ϕ_j is the contribution of feature j to the individual prediction. SHAP values satisfy three key properties: *local accuracy*, *missingness*, and *consistency*.

5.2 SHAP Implementation

```
1 import shap
2
3 # Create SHAP explainer for LightGBM
4 explainer = shap.TreeExplainer(lightgbm_model)
5
6 # Compute SHAP values for test set
7 shap_values = explainer.shap_values(X_test)
8
9 # Per-prediction explanation
10 shap.force_plot(
11     explainer.expected_value,
12     shap_values[0],
13     X_test.iloc[0],
14     feature_names=feature_names
15 )
16
17 # Global summary plot (feature importance)
18 shap.summary_plot(shap_values, X_test, feature_names=feature_names)
19
20 # Waterfall chart for single prediction
21 shap.waterfall_plot(
22     shap.Explanation(
23         values=shap_values[0],
24         base_values=explainer.expected_value,
25         data=X_test.iloc[0]
26     )
27 )
```

Listing 1: SHAP Integration with LightGBM

5.3 SHAP Visualisation Types

1. **Summary Plot:** Shows global feature importance ranked by mean absolute SHAP value, with colour indicating feature magnitude. Confirms `model`, `engine_cc`, and `mileage_per_year` as top drivers.
2. **Force Plot (per prediction):** Visualises how each feature pushes the prediction above or below the base value (mean predicted price). Features in red increase the prediction; blue features decrease it.
3. **Waterfall Chart:** A cumulative view of feature contributions for a single prediction, showing the step-by-step build-up from the base value to the final prediction.
4. **Dependence Plot:** Shows the relationship between a single feature and its SHAP value across the dataset, revealing non-linear effects and interaction patterns.

5.4 Example Explanation

For a 2020 Toyota Corolla (Local, 1800cc, 45,000 km, Lahore):

- Base value (mean prediction): Rs 4,120,000
- `model` (Corolla, premium variant): +Rs 380,000

- **engine_cc** (1800cc, above median): +Rs 210,000
- **mileage_per_year** (low annual usage): +Rs 95,000
- **car_age** (6 years): −Rs 255,000
- **Final prediction:** Rs 4,550,000

5.5 XAI Business Value

- **Buyers:** Understand exactly what features are driving the price up or down.
- **Sellers:** Know which improvements (e.g. servicing to reduce mileage perception) can increase resale value.
- **Financial Institutions:** Obtain auditable, feature-level price justifications for loan collateral.
- **Regulators:** Verify that pricing recommendations are free of discriminatory patterns.

6 Phase 6: Deployment

6.1 Streamlit Web Application

The final system is deployed as a fully functional web application accessible at:

<https://muhammad-ahmad2511-used-car-price-predictor.streamlit.app>

6.2 Application Features

1. **Real-Time Price Prediction:** Users input car specifications (make, model, year, mileage, engine, assembly, city, transmission, fuel type) and receive an instant price estimate in PKR.
2. **SHAP Explanation Panel:** Every prediction is accompanied by a waterfall or force plot showing the contribution of each feature to the predicted price.
3. **Interactive Inputs:** Drop-down menus populated with actual values from the training dataset ensure valid inputs at all times.
4. **Price Range Confidence:** Along with a point estimate, a confidence interval based on model residuals is displayed.
5. **Market Comparison:** The predicted price is contextualised against similar listings in the dataset.

6.3 Technology Stack

Table 14: Technology Stack

Category	Tools
Language	Python 3.11
Web Scraping	BeautifulSoup4, Requests
Data Processing	Pandas $\geq 2.2.0$, NumPy $\geq 2.0.0$
Machine Learning	XGBoost $\geq 2.1.0$, LightGBM $\geq 4.3.0$
Explainability	SHAP $\geq 0.45.0$
Visualisation	Matplotlib $\geq 3.8.0$, Seaborn $\geq 0.13.0$
Frontend	Streamlit
Deployment	Streamlit Cloud
Version Control	Git, GitHub

6.4 Deployment Architecture

1. **Model Serialisation:** Trained LightGBM model and label encoders serialised as .pkl files using `joblib`.
2. **Cloud Hosting:** Application hosted on Streamlit Community Cloud, pulling from the GitHub repository on each deployment trigger.
3. **Dependency Management:** `requirements.txt` specifies all Python packages with version floor constraints for Python 3.14 compatibility.
4. **Inference Pipeline:** User inputs $\rightarrow$ feature engineering $\rightarrow$ label encoding $\rightarrow$ LightGBM inference $\rightarrow$ SHAP explanation $\rightarrow$ display.

7 Team Contributions

Table 15: Individual Contributions by Module

Team Member	GitHub Handle	Modules & Contributions
Muhammad Ahmad (23L-2511)	Muhammad-Ahmad2511	XGBoost model implementation, hyperparameter tuning (GridSearchCV), model serialisation (.pkl files), initial project setup, Streamlit app frontend
Muhammad Sufyan (23L-2518)	s6090-8798	LightGBM model implementation, hyperparameter tuning, model weights management, README documentation, app backend integration
Muhammad Daniyal (23L-2600)	KingofMagician	Model evaluation script, performance comparison analysis, SHAP integration, visualisation generation, requirements.txt setup, report writing

Table 16: Commit Distribution (All Deliverables)

Team Member	Commits	Lines Added	Files Modified
Muhammad Ahmad (23L-2511)	3+	~550	5+
Muhammad Sufyan (23L-2518)	3+	~480	4+
Muhammad Daniyal (23L-2600)	3+	~420	6+
Total	9+	~1,450	15+

8 Conclusion

8.1 Summary of Achievements

This project successfully delivered a complete AI pipeline addressing price opacity in Pakistan’s used car market:

1. **Data Pipeline:** Scraped 24,728 records from PakWheels.com; 97.6% retention after a 7-step cleaning process; 15 engineered features capturing domain-specific patterns.
2. **Model Performance:** LightGBM achieves $R^2=0.873$ and $MAE=Rs\ 562K$ on the test set — state-of-the-art performance for this domain.
3. **Explainability:** SHAP integration provides per-prediction transparency, fulfilling the XAI requirement and building user trust.

4. **Deployment:** Fully operational Streamlit application deployed on Streamlit Cloud with real-time inference.
5. **Fairness:** No significant bias detected across city tiers or assembly types.

8.2 Limitations

- Performance degrades for the luxury segment (>Rs 10M) due to data scarcity.
- Model requires periodic retraining as market prices shift.
- Scraped data reflects listed (not transacted) prices, which may include negotiation premiums.

8.3 Future Work

- **Neural approaches:** Tabular deep learning (TabNet, FT-Transformer) for improved luxury-segment accuracy.
- **Temporal modelling:** Incorporate price time-series to capture market trends and seasonal effects.
- **Image features:** Integrate vehicle photo analysis (condition assessment) as additional features.
- **API endpoint:** Expose the model as a REST API for integration with dealer management systems.
- **Continuous retraining:** Automated pipeline to retrain the model monthly with fresh scraped data.

References

1. PakWheels.com — Pakistan’s Largest Automotive Marketplace. <https://www.pakwheels.com/used-cars/search/-/>
2. Chen, T., & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. *KDD '16*.
3. Ke, G., et al. (2017). LightGBM: A Highly Efficient Gradient Boosting Decision Tree. *NeurIPS 2017*.
4. Lundberg, S. M., & Lee, S.-I. (2017). A Unified Approach to Interpreting Model Predictions. *NeurIPS 2017*.
5. Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. *JMLR*, 12, 2825–2830.